

SQL essentials

Mastering database queries:
a comprehensive guide to SQL

Learning session 06
Using set operators to combine
the results of multiple queries

Instructor

Péter Fülöp

peteristvan.fulop@hcl.com



Introduction to set theory

The application of set theory in SQL

- **Set Theory** is a mathematical concept where a set is defined as a collection of elements. This collection adheres to **specific rules**, including the **absence of duplicates**, and **the elements themselves are not arranged** in any particular order.
- The basic principle of set theory is that **an element belongs to a set**.
- In SQL, **the outcome of a query can be viewed as a set**, illustrating the application of set theory.
- When **applying set operators in SQL**, the SQL sets involved must adhere to two criteria: they should contain an equal number of columns, and the corresponding columns in each set must share the same datatype.

Set operators

UNION

This operator combines the result sets of two or more SELECT statements and eliminates duplicate rows. The result is a set that contains all distinct rows from the queries involved.

INTERSECT

The **INTERSECT set operator** is used to return the intersection of two or more queries. Essentially, it finds the common elements that exist in the result sets of both queries. **The number and order of columns** in the SELECT statements must be the same. **The data types of the corresponding columns** must be compatible or, ideally, the same.

EXCEPT

The **EXCEPT set operator** is used to return all rows from the first (left) query that are not found in the second (right) query. It effectively subtracts the result of one query from another.

Example 1: using the UNION set operator



The company aims to assign team leads for employees who currently do not have one. To facilitate this, a list should be created that includes both employees lacking a team lead and those who are already serving as team leads. Additionally, the list should feature a column titled 'is_team_lead', marked with 'Y' or 'N' to differentiate between the roles.

Let's split our approach into **two distinct queries**, subsequently merging their results. **The first query** will focus on employees without a team lead, while the **second query** will compose a list of existing team leads.

IMPORTANT: both queries should return the same columns in the same order.

SQL

```
query1
UNION
query2
```

The UNION set operator

Example 1: Using the UNION set operator



(Part 1) Create a list that includes all the employees lacking a team lead.

SQL

```
SELECT id, first_name, last_name, 'N' as is_team_lead
FROM employees
WHERE lead_employee_id IS NULL;
```

OUTPUT

	<input type="checkbox"/> id	<input type="checkbox"/> first_name	<input type="checkbox"/> last_name	<input type="checkbox"/> is_team_lead
1	16	Eva	Kiss	N
2	17	Tamas	Molnar	N
3	18	Zsuzsa	Papp	N
4	19	Laszlo	Varga	N
5	20	Noemi	Biro	N
6	10	Laszlo	Nagy	N

6 rows

The UNION set operator

Example 1: Using the UNION set operator



(Part 2) Create a list that includes all team leads.

```
SQL

SELECT emp.id, emp.first_name, emp.last_name, 'Y' as is_team_lead
FROM employees emp
    INNER JOIN (SELECT DISTINCT lead_employee_id FROM employees) AS tld
        ON emp.id = tld.lead_employee_id;
```

OUTPUT

	<input type="checkbox"/> id	<input type="checkbox"/> first_name	<input type="checkbox"/> last_name	<input type="checkbox"/> is_team_lead
1	1	Gabor	Nagy	Y
2	2	Istvan	Kiss	Y
3	3	Zoltan	Toth	Y
4	4	Eva	Horvath	Y

10 rows

The UNION set operator

Example 1: Using the UNION set operator

Let's merge the results of the two queries using the UNION set operator.

```
SQL

(SELECT id, first_name, last_name, 'N' as is_team_lead
FROM employees
WHERE lead_employee_id IS NULL)

UNION

(SELECT emp.id, emp.first_name, emp.last_name, 'Y' as is_team_lead
FROM employees emp
INNER JOIN (SELECT DISTINCT lead_employee_id FROM employees) AS tld
ON emp.id = tld.lead_employee_id)
ORDER BY last_name ASC, first_name ASC;
```

Example 2: Using the UNION set operator



The firm resolved to implement a salary hike of 10% for employees engaged in active projects, and a 5% raise for those not involved in active projects but who had their most recent salary review six months ago or more. The list should include employee ID, first names, last names, present salary, percentage of salary increment, and the amount of the raise. This list should be arranged alphabetically by last names and then by first names.

Let's break down our solution into **two separate queries** and then combine or merge the resulting rows.

The first query will determine those employees who are going to receive 10% salary raise and **the second query** will determine employees receiving 5% salary raise. Both queries will return the same columns in the same order.

A screenshot of a Mac OS X-style application window titled "SQL". The window contains two lines of text: "query1" on the first line and "UNION" on the second line, followed by "query2" on the third line. The window has the standard red, yellow, and green close buttons in the top-left corner.

Example 2: Using the UNION set operator

The first query below solves the following section of the previous task:

-  **(PART 1) The firm resolved to implement a salary hike of 10% for employees engaged in active projects.**
The list should include employee ID, first names, last names, present salary, percentage of salary increment, and the amount of the raise.

● ● ● QUERY 1 SQL

```
SELECT emp.id, emp.first_name, emp.last_name, emp.salary,
       10 as salary_raise_percent, 0.1*salary as salary_raise
  FROM employees emp
 INNER JOIN project_assignments as pra ON emp.id = pra.employee_id
 INNER JOIN projects as p ON pra.project_id = p.id
 WHERE COALESCE(p.start_date, current_date + 1) <= current_date
       AND p.completion_date IS NULL
       AND p.cancellation_date IS NULL
       AND COALESCE(pra.start_date, current_date + 1) <= current_date
       AND pra.end_date IS NULL;
```

The UNION set operator

Example 2: Using the UNION set operator

The second query below solves the following section of the previous task:

- (PART 2)** The firm resolved to implement a 5% raise for those not involved in active projects but who had their most recent salary review six months ago or more. The list should include employee ID, first names, last names, present salary, percentage of salary increment, and the amount of the raise.

A horizontal bar featuring three large, solid-colored circles on the left side: red, yellow, and green. To the right of the circles, the word "QUERY" is written in a large, dark gray, sans-serif font, followed by the number "2" in a slightly smaller size.

SQL

```
SELECT emp.id, emp.first_name, emp.last_name, emp.salary, 5 as salary_raise_percent, 0.05 * salary
FROM employees emp
WHERE COALESCE(current_date - emp.last_salary_review_date, 0) >= 180
AND emp.id NOT IN (SELECT emp1.id
                     FROM employees emp1
                     INNER JOIN project_assignments as pra1 ON emp1.id = pra1.employee_id
                     INNER JOIN projects as p1 ON pra1.project_id = p1.id
                     WHERE COALESCE(p1.start_date, current_date + 1) <= current_date
                     AND p1.completion_date IS NULL
                     AND p1.cancellation_date IS NULL
                     AND COALESCE(pra1.start_date, current_date + 1) <= current_date
                     AND pra1.end_date IS NULL)
```

The UNION set operator

Example 2: Using the UNION set operator

Combining the results of the queries shown on the previous two pages you should get the results below:

OUTPUT

	<input type="checkbox"/> id	<input type="checkbox"/> first_name	<input type="checkbox"/> last_name	<input type="checkbox"/> salary	<input type="checkbox"/> salary_raise_percent	<input type="checkbox"/> salary_raise
1	27	Andras	Farkas	3500	5	175
2	4	Eva	Horvath	4500	10	450
3	15	Istvan	Horvath	4500	5	225
4	23	Ferenc	Kiss	2900	5	145
5	2	Istvan	Kiss	3500	10	350
6	6	Andras	Kovacs	2600	10	260
7	11	Anna	Kovacs	3500	10	350
8	7	Anita	Molnar	3800	10	380
9	22	Anna	Nagy	3800	5	190
10	29	Bela	Nagy	3100	5	155
11	1	Gabor	Nagy	2500	10	250
12	10	Laszlo	Nagy	3200	10	320
13	19	Zsuzsanna	Papp	2800	5	140

18 rows

Example 2: refactoring - increasing code readability

The queries shown in example 2 can be restructured using Common Table Expressions (CTEs), which will greatly enhance the readability of the code.

QUERY REFACTORING (PART 1) SQL

```
WITH employees_on_active_project AS
  (SELECT emp1.id
   FROM employees emp1
        INNER JOIN project_assignments as pra1
          ON emp1.id = pra1.employee_id
        INNER JOIN projects as p1 ON pra1.project_id = p1.id
   WHERE COALESCE(p1.start_date, current_date + 1) <= current_date
     AND p1.completion_date IS NULL
     AND p1.cancellation_date IS NULL
     AND COALESCE(pra1.start_date, current_date + 1) <= current_date
     AND pra1.end_date IS NULL )
```

Example 2: refactoring - increasing code readability

Let us use the defined CTE in refactoring our queries.

QUERY REFACTORYING (PART 2)

SQL

```
WITH employees_on_active_project AS (...)

SELECT emp.id, emp.first_name, emp.last_name, emp.salary,
       10 as salary_raise_percent, 0.1 * salary as salary_raise
FROM employees emp
      INNER JOIN employees_on_active_project eoap ON emp.id = eoap.id
UNION
SELECT emp.id, emp.first_name, emp.last_name, emp.salary,
       5 as salary_raise_percent, 0.05 * salary
FROM employees emp
      LEFT OUTER JOIN employees_on_active_project eoap ON emp.id = eoap.id
WHERE COALESCE(current_date - emp.last_salary_review_date, 0) >= 180
      AND eoap.id IS NULL
ORDER BY last_name, first_name;
```

The UNION set operator

Retaining/removing duplicates

The UNION set operator eliminates duplicate rows; to retain duplicates, use UNION ALL instead.

SQL

```
select * from employees
UNION ALL
select * from employees
order by id;
```

OUTPUT

	<input type="checkbox"/> id	<input type="checkbox"/> first_name	<input type="checkbox"/> middle_name	<input type="checkbox"/> last_name	<input type="checkbox"/> hire_date	<input type="checkbox"/> b
1	1	Gabor	<null>	Nagy	2022-02-15	199
2	1	Gabor	<null>	Nagy	2022-02-15	199
3	2	Istvan	Peter	Kiss	2022-03-01	198
4	2	Istvan	Peter	Kiss	2022-03-01	198

58 rows

The INTERSECT set operator

Using the INTERSECT set operator



Which languages are spoken in both Argentina and Chile?

Do not include generic terms such as 'indigenous', 'Other', or 'unspecified'. Provide the names of specific languages.

Order the list of languages in the alphabetical order.

Let's break down our solution into **two separate queries** and then intersect the resulting rows.

The first query will determine languages that are spoken in Argentina. The second query will determine languages spoken in Chile.

A screenshot of a computer screen showing a SQL editor window. The window has a title bar with three colored buttons (red, yellow, green) on the left and the word "SQL" on the right. The main area of the window contains the following text:

```
query1  
INTERSECT  
query2
```

The INTERSECT set operator

Using the INTERSECT set operator



(PART 1) Which languages are spoken in Argentina?

Do not include generic terms such as 'indigenous', 'Other', or 'unspecified'. Provide the names of specific languages.

SQL

```
SELECT name AS language
FROM languages
WHERE languages.code = (select code from countries where lower(name) = 'argentina')
    AND lower(languages.name) NOT IN ('indigenous', 'other', 'unspecified');
```

OUTPUT

	language
1	Spanish
2	Italian
3	English
4	German
5	French

5 rows

The INTERSECT set operator

Using the INTERSECT set operator



(PART 2) Which languages are spoken in Chile?

Do not include generic terms such as 'indigenous', 'Other', or 'unspecified'. Provide the names of specific languages.

SQL

```
SELECT name AS language
FROM languages
WHERE languages.code = (select code from countries where lower(name) = 'chile')
    AND lower(languages.name) NOT IN ('indigenous', 'other', 'unspecified');
```

OUTPUT

	language
1	Spanish
2	English

2 rows

The INTERSECT set operator

Using the INTERSECT set operator

Let's intersect the resulting rows using the INTERSECT operator.

SQL

```
SELECT name AS language
FROM languages
WHERE languages.code = (select code from countries where lower(name) = 'argentina')
    AND lower(languages.name) NOT IN ('indigenous', 'other', 'unspecified')
INTERSECT
SELECT name AS language
FROM languages
WHERE languages.code = (select code from countries where lower(name) = 'chile')
    AND lower(languages.name) NOT IN ('indigenous', 'other', 'unspecified')
ORDER BY language ASC;
```

OUTPUT

	language
1	English
2	Spanish

2 rows

Using the EXCEPT set operator



Find the languages that are not the official language of any country.
Order the list of languages in the alphabetical order.

We can approach our solution in **two steps** by executing two separate queries and then combining the results using the EXCEPT set operator.

In the **first query**, we will identify languages that are spoken but are not recognized as official languages in at least one country. In the **second query**, we will specifically retrieve the official languages.

A screenshot of a SQL editor window. The title bar has three colored dots (red, yellow, green) on the left and the word "SQL" on the right. The main area contains the following SQL code:

```
query1  
EXCEPT  
query2
```

The EXCEPT set operator

Using the EXCEPT set operator



(PART 1) Which languages are official languages in at least one country.



SQL

```
SELECT name as language  
FROM languages  
WHERE official IS TRUE;
```



OUTPUT

	language
1	Turkic
2	Other
3	Greek
4	Other
5	unspecified
6	French

629 rows

Using the EXCEPT set operator



(PART 2) Which languages are spoken but are not recognized as official languages in at least one country.

SQL

```
SELECT name as language
FROM languages
WHERE official IS NOT TRUE;
```

OUTPUT

	language
1	Dari
2	Pashto
3	Albanian
4	Arabic
5	Berber or Tamazight
6	Catalan
7	Portuguese

294 rows

The EXCEPT set operator

Using the EXCEPT set operator

Let's combine the resulting rows using the EXCEPT operator.

SQL

```
SELECT name AS language
FROM languages
WHERE official IS NOT TRUE
EXCEPT
SELECT name AS language
FROM languages
WHERE official IS TRUE
ORDER BY language ASC;
```

OUTPUT

	language
1	Akyem
2	Alsatian
3	Amerindian

253 rows